

```
1 import pandas as pd
2 import numpy as np
3 import time
```

```
1 from google.colab import drive
2 drive.mount('/content/drive/')

```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount(force=True)

Ingesting The Data

```
1 train_data = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/benz/train.csv')
2 test_data = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/data/benz/test.csv')
```

Data Analysis

```
1 train_data.shape
```

```
(4209, 378)
```

```
1 test_data.shape
```

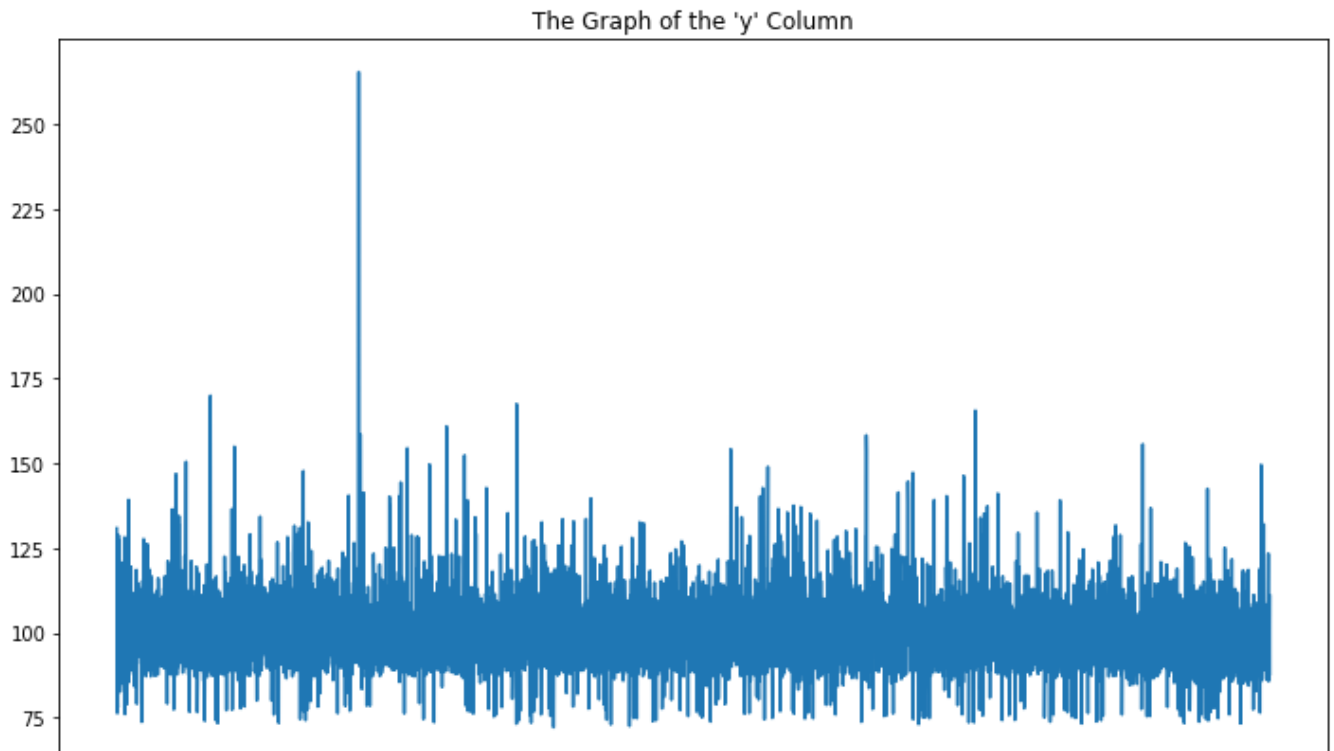
```
(4209, 377)
```

```
1 train_data.head()
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	

5 rows × 378 columns

```
1 train_data['y'].plot(figsize=(12,7), title="The Graph of the 'y' Column");
```



```
1 test_data.head()
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X381
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	(
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	(
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	(
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	(
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	(

5 rows × 377 columns




```
1 train_data.dtypes
```

```
ID          int64
y          float64
X0          object
X1          object
X2          object
...
X380        int64
X382        int64
X383        int64
X384        int64
X385        int64
Length: 378, dtype: object
```

```
1 train_data.describe()
```

	ID	y	X10	X11	X12	X13	X14
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	4209.000000	4209.000000
mean	4205.960798	100.669318	0.013305	0.0	0.075077	0.057971	0.428131
std	2437.608688	12.679381	0.114590	0.0	0.263547	0.233716	0.494861
min	0.000000	72.110000	0.000000	0.0	0.000000	0.000000	0.000000
25%	2095.000000	90.820000	0.000000	0.0	0.000000	0.000000	0.000000
50%	4220.000000	99.150000	0.000000	0.0	0.000000	0.000000	0.000000
75%	6314.000000	109.010000	0.000000	0.0	0.000000	0.000000	1.000000
max	8417.000000	265.320000	1.000000	0.0	1.000000	1.000000	1.000000

8 rows × 370 columns



1 test_data.dtypes

```
ID      int64
X0      object
X1      object
X2      object
X3      object
...
X380    int64
X382    int64
X383    int64
X384    int64
X385    int64
Length: 377, dtype: object
```

1 test_data.describe()

	ID	X10	X11	X12	X13	X14	
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000

Checking Null Values

std	2423.070920	0.130303	0.013414	0.202334	0.233400	0.434032	0.130303
-----	-------------	----------	----------	----------	----------	----------	----------

```
1 train_data.isnull().sum()
```

ID	0
y	0
X0	0
X1	0
X2	0
..	
X380	0
X382	0
X383	0
X384	0
X385	0
Length: 378, dtype: int64	

```
1 test_data.isnull().sum()
```

ID	0
X0	0
X1	0
X2	0
X3	0
..	
X380	0
X382	0
X383	0
X384	0
X385	0
Length: 377, dtype: int64	

```
1 train_data.describe(include='object')
```

	X0	X1	X2	X3	X4	X5	X6	X8	
count	4209	4209	4209	4209	4209	4209	4209	4209	
unique	47	27	44	7	4	29	12	25	
top	z	aa	as	c	d	w	g	j	
freq	360	833	1659	1942	4205	231	1042	277	

Separating categorical and numeric feature

```

1 dictionary={}
2 dictionary['num'] = train_data.dtypes[train_data.dtypes=='int64'].index
3 dictionary['cat'] = train_data.dtypes[train_data.dtypes=='object'].index
4 dictionary

{'cat': Index(['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8'], dtype='object'),
 'num': Index(['ID', 'X10', 'X11', 'X12', 'X13', 'X14', 'X15', 'X16', 'X17', 'X18',
              ...
              'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
              'X385'],
              dtype='object', length=369)}

```

```
1 train_data['X376'].value_counts()
```

```

0    3968
1     241
Name: X376, dtype: int64

```

Finding and removing features with zero variance

```

1 da=[]
2 count=0
3 for i in dictionary['num']:
4     if(np.var(train_data[i])==0):
5         da.append(i)
6 print(da)

['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330',

```

```

1 z=['ID','y']
2 da.extend(z)
3 print(da)

['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330',

```

```

1 y_train=train_data['y'].values
2 ID_train=train_data['ID'].values
3 ID_test=test_data['ID'].values

```

```
1 da_test=['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', '>
```

```

1 X_train=train_data.drop(da,axis=1)
2 X_test=test_data.drop(da_test,axis=1)

```

```
1 X_train.shape
```

```
(4209, 364)
```

```
1 X_test.shape
```

```
(4209, 364)
```

Applying Label encoding on the categorical variables

```
1 from sklearn.preprocessing import LabelEncoder
```

```
1 #Train dataset
```

```
2 le0=LabelEncoder()
```

```
3 le1=LabelEncoder()
```

```
4 le2=LabelEncoder()
```

```
5 le3=LabelEncoder()
```

```
6 le4=LabelEncoder()
```

```
7 le5=LabelEncoder()
```

```
8 le6=LabelEncoder()
```

```
9 le8=LabelEncoder()
```

```
1 le0.fit(X_train['X0'])
```

```
2 le1.fit(X_train['X1'])
```

```
3 le2.fit(X_train['X2'])
```

```
4 le3.fit(X_train['X3'])
```

```
5 le4.fit(X_train['X4'])
```

```
6 le5.fit(X_train['X5'])
```

```
7 le6.fit(X_train['X6'])
```

```
8 le8.fit(X_train['X8'])
```

```
LabelEncoder()
```

```
1 #Test Dataset
```

```
2 let0=LabelEncoder()
```

```
3 let1=LabelEncoder()
```

```
4 let2=LabelEncoder()
```

```
5 let3=LabelEncoder()
```

```
6 let4=LabelEncoder()
```

```
7 let5=LabelEncoder()
```

```
8 let6=LabelEncoder()
```

```
9 let8=LabelEncoder()
```

```
1 #Test Dataset
```

```
2 let0.fit(X_test['X0'])
```

```
3 let1.fit(X_test['X1'])
```

```
4 let2.fit(X_test['X2'])
```

```
5 let3.fit(X_test['X3'])
```

```

6 let4.fit(X_test['X4'])
7 let5.fit(X_test['X5'])
8 let6.fit(X_test['X6'])
9 let8.fit(X_test['X8'])

```

```
LabelEncoder()
```

Transforming and replacing the categorical variables into 0s and 1s

```

1 #Train dataset
2 X_train['X0'] = le0.transform(X_train['X0'])
3 X_train['X1'] = le1.transform(X_train['X1'])
4 X_train['X2'] = le2.transform(X_train['X2'])
5 X_train['X3'] = le3.transform(X_train['X3'])
6 X_train['X4'] = le4.transform(X_train['X4'])
7 X_train['X5'] = le5.transform(X_train['X5'])
8 X_train['X6'] = le6.transform(X_train['X6'])
9 X_train['X8'] = le8.transform(X_train['X8'])

```

```

1 #Test dataset
2 X_test['X0'] = let0.transform(X_test['X0'])
3 X_test['X1'] = let1.transform(X_test['X1'])
4 X_test['X2'] = let2.transform(X_test['X2'])
5 X_test['X3'] = let3.transform(X_test['X3'])
6 X_test['X4'] = let4.transform(X_test['X4'])
7 X_test['X5'] = let5.transform(X_test['X5'])
8 X_test['X6'] = let6.transform(X_test['X6'])
9 X_test['X8'] = let8.transform(X_test['X8'])

```

```
1 X_train.head()
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	X380	X381
0	32	23	17	0	3	24	9	14	0	0	...	0	0	1	0	0	0	0
1	32	21	19	4	3	28	11	14	0	0	...	1	0	0	0	0	0	0
2	20	24	34	2	3	27	9	23	0	0	...	0	0	0	0	0	0	0
3	20	21	34	5	3	27	11	4	0	0	...	0	0	0	0	0	0	0
4	20	23	34	5	3	12	3	13	0	0	...	0	0	0	0	0	0	0

5 rows × 364 columns

```
1 X_test.head()
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	X379	X380	X381
0	21	23	34	5	3	26	0	22	0	0	...	0	0	0	1	0	0	
1	42	3	8	0	3	9	6	24	0	0	...	0	0	1	0	0	0	
2	21	23	17	5	3	0	9	9	0	0	...	0	0	0	1	0	0	
3	21	13	34	5	3	31	11	13	0	0	...	0	0	0	1	0	0	
4	45	20	17	2	3	30	8	12	0	0	...	1	0	0	0	0	0	

Perform dimensionality reduction on Train dataset

```
1 from sklearn.decomposition import PCA
2 pca = PCA(n_components=0.7)
3 pca.fit(X_train)
```

```
PCA(n_components=0.7)
```

```
1 pca.explained_variance_ratio_
```

```
array([0.38334782, 0.21388033, 0.13261866])
```

```
1 X_train_transformed = pca.transform(X_train)
2 X_test_transformed = pca.transform(X_test)
```

```
1 X_train_transformed.shape, X_test_transformed.shape
```

```
((4209, 3), (4209, 3))
```

```
1 pca.inverse_transform(X_train_transformed).shape
```

```
(4209, 364)
```

Split the datasets

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X_train_transformed, y_train, random_
```

```
1 print(X_train.shape)
2 print(X_test.shape)
3 print(y_train.shape)
4 print(y_test.shape)
```

```
(3156, 3)
(1053, 3)
```



```
(3156,)
(1053,)
```

Evaluate with ensemble learning algorithms

```
1 import xgboost
2 from sklearn.ensemble import GradientBoostingRegressor
3 from sklearn.metrics import mean_squared_error

1 #Using XGBoost - The most powerful ensemble learning algorithm
2 start = time.time()
3 xgbr = xgboost.XGBRegressor()
4 xgbr.fit(X_train, y_train)
5
6 y_pred= xgbr.predict(X_test)
7 score = mean_squared_error(y_pred, y_test)
8 print('Score : ', score)
9 end = time.time()
10
11 print('\nRunTime : ', end - start)
```

```
[02:53:42] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated
Score : 110.00614368130348
```

```
RunTime : 0.1572103500366211
```



```
1 #Using GradientBoost
2 start = time.time()
3 gbr = GradientBoostingRegressor()
4 gbr.fit(X_train, y_train)
5
6 y_pred= gbr.predict(X_test)
7 score = mean_squared_error(y_pred, y_test)
8 print('Score : ', score)
9 end = time.time()
10
11 print('\nRunTime : ', end - start)
```

```
Score : 111.54442824296098
```

```
RunTime : 0.3748013973236084
```

As seen above, XGBoost performs more efficiently than GradientBoost

✓ 0s completed at 10:53 PM

● ✕